

Hive Case Study

E-Commerce Sales Review

Problem Statement: The tech companies are exploring different ways to improve their sales by analysing customer behaviour and gaining insights about product trends, due to the increasing popularity in online sales. Therefore, as a big data analyst, we are expected to extract data and gather insights from a real-life data set of an e-commerce company.

Objective: Need to analyse and gain insights about the clickstream data from a website so that we can extract insights about the customers behaviour.

The steps involved in the entire process are as follows:

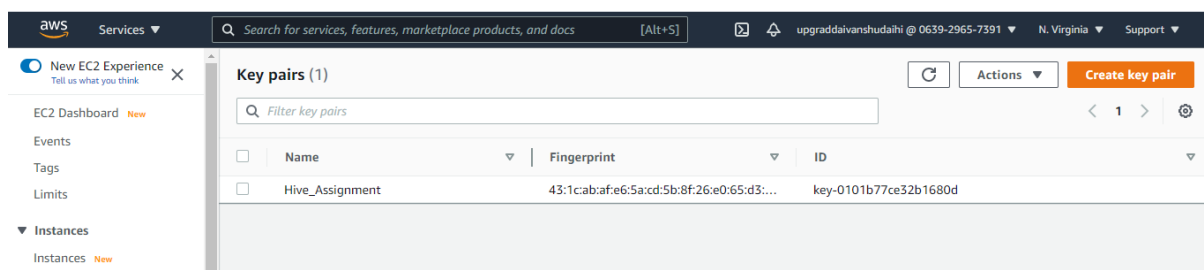
- Copying the data set into the HDFS:
 - Launch an EMR cluster that utilizes the Hive services, and
 - Move the data from the S3 bucket into the HDFS
- Creating the database and launching Hive queries on your EMR cluster:
 - Create the structure of your database,
 - Use optimized techniques to run your queries as efficiently as possible
 - Show the improvement of the performance after using optimization on any single query.
 - Run Hive queries to answer the questions given below.

Launch the EMR Cluster.

In order to launch an EMR Cluster, the following steps were followed:

- Create a key-pair and download the PEM/ PPK file.

Creation of Key Pair:



- After creation of Key-pair, now we need create an EMR cluster. While creating an EMR cluster need to make sure that we are selecting **m4.large** Master and Core node of single instances. Also need to select the correct key-pair in the security option

Node type	Instance type	Instance count	Purchasing option
Master Master - 1	m4.large 2 vCore, 8 GiB memory, EBS only storage EBS Storage: 32 GiB Add configuration settings	1 Instances	On-demand Spot Use on-demand as max price
Core Core - 2	m4.large 2 vCore, 8 GiB memory, EBS only storage EBS Storage: 32 GiB Add configuration settings	1 Instances	On-demand Spot Use on-demand as max price

Step 1: Software and Steps
Step 2: Hardware
Step 3: General Cluster Settings
Step 4: Security

General Options

Cluster name: Hive_Case_Study

Logging: ☒
S3 folder: s3://aws-logs-063929657391-us-east-1/elasticmapreduce
Debugging: ☒
Termination protection: ☒

Tags

Key	Value (optional)
Add a key to create a tag	

Additional Options

EMRFS consistent view: ☐
Custom AMI ID: None

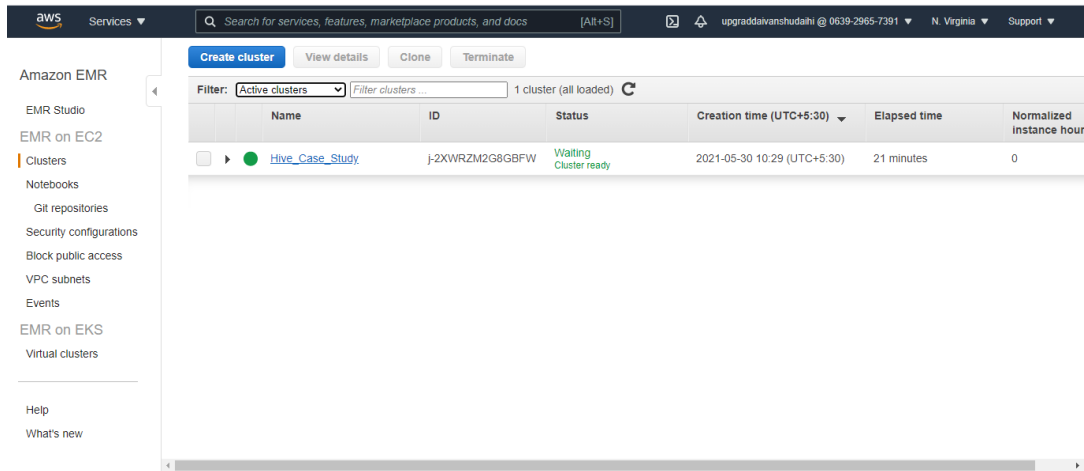
Step 1: Software and Steps
Step 2: Hardware
Step 3: General Cluster Settings
Step 4: Security

Security Options

EC2 key pair: Hive_Assignment
Cluster visible to all IAM users in account: ☒
Permissions: Default
EMR role: EMR_DefaultRole
EC2 Instance profile: EMR_EC2_DefaultRole
Auto Scaling role: EMR_AutoScaling_DefaultRole
Security Configuration
EC2 security groups

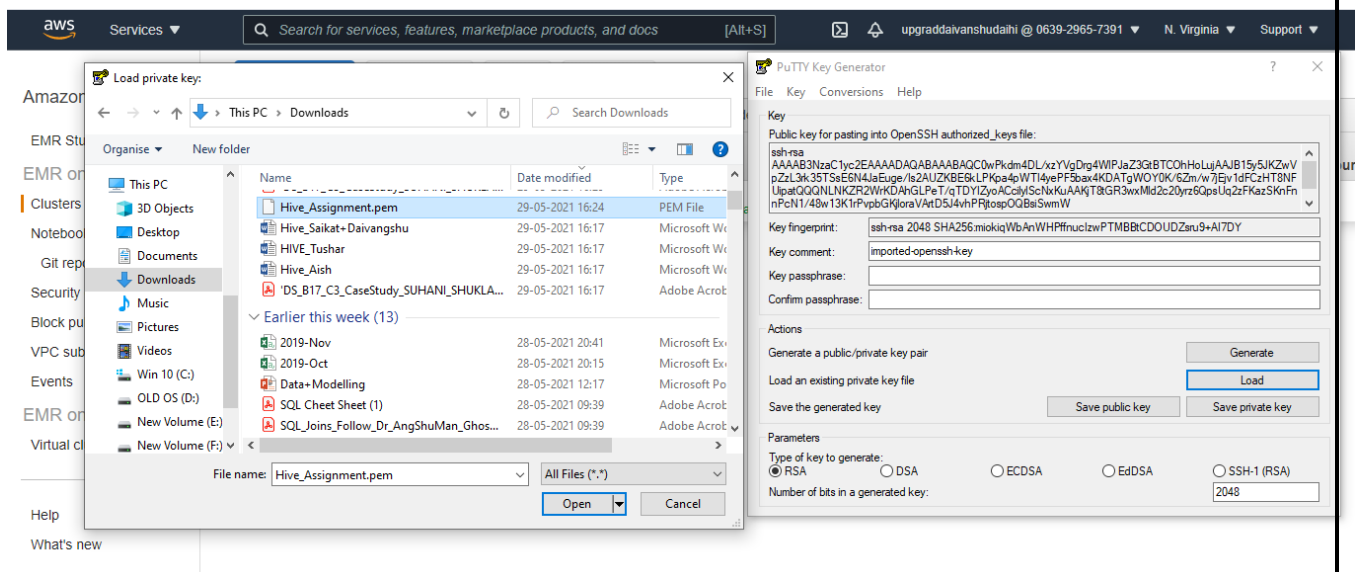
Cancel Previous Create cluster

- Now the cluster is created and we are ready to move to the next stage i.e moving data from S3 to HDFS.

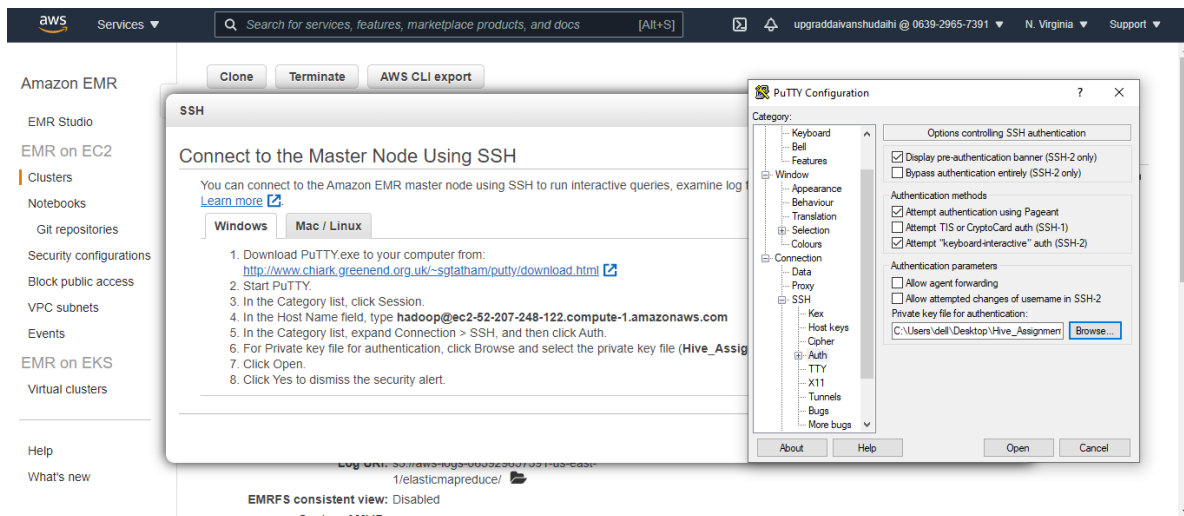


Connecting to EMR cluster:

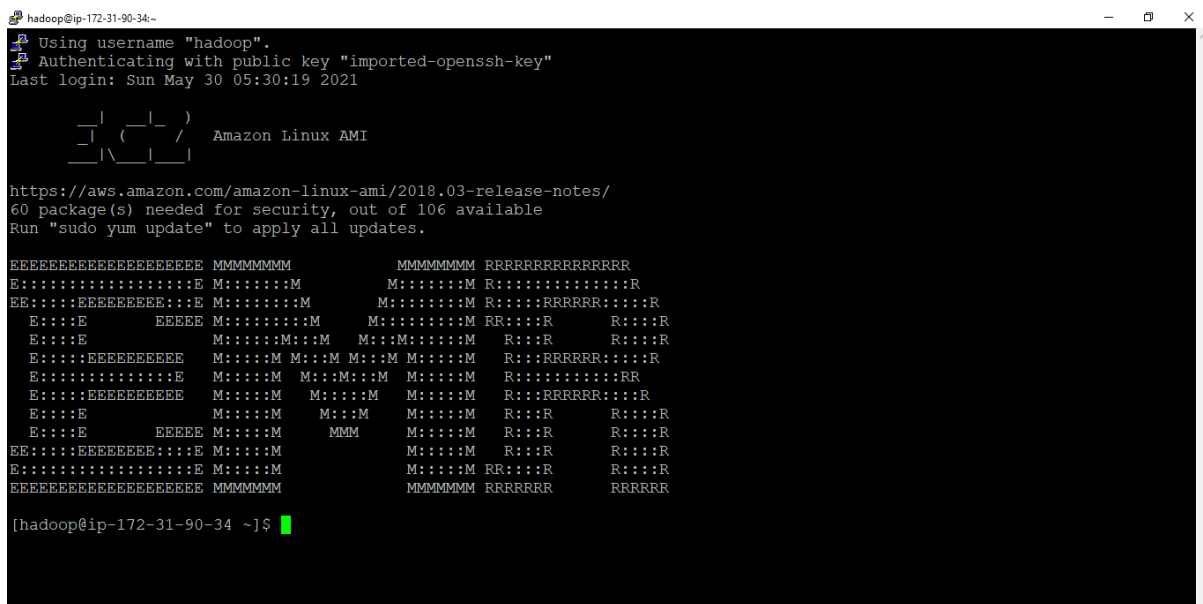
- Here need to open PuTTYgen application for windows machine and we have to Load the previously downloaded .pem Key-pair file and save the private key which is in the extension .ppk



Once cluster in running state we have to click on Master public DNS. We have to open the putty configuration and then give the host name (master node DNS) and then browse to the private key file location by clicking on Connection → SSH → Auth. Now we need to open Putty and connect to the master node by selecting the .ppk file.




Connection to hadoop is successful:



Load the data sets into HDFS from S3:

- Create a directory named 'Hive_assignment' in Hadoop.



- Move the data from the s3 buckets to the HDFS using the distributed copy command.

Loading the s3 public data set to created directory "Hive_assignment" in hadoop .

Command: **hadoop distcp 's3://e-commerce-events-ml/2019-Oct.csv' / Hive_assignment /2019-Oct.csv**

hadoop distcp 's3://e-commerce-events-ml/2019-Nov.csv'/Hive_assignment/2019-Nov.csv

```
[hadoop@ip-172-31-91-114 ~]$ hadoop distcp s3://e-commerce-events-ml/2019-Oct.csv /Hive_assignment/2019-Oct.csv
```

```
DistCp Counters
  Bytes Copied=482542278
  Bytes Expected=482542278
  Files Copied=1
```

```
[hadoop@ip-172-31-90-34 ~]$ hadoop distcp s3://e-commerce-events-ml/2019-Nov.csv /Hive_assignment/2019-Nov.csv
```

```
DistCp Counters
  Bytes Copied=545839412
  Bytes Expected=545839412
  Files Copied=1
```

- View the data in HDFS by executing below commands

```
[hadoop@ip-172-31-91-114 ~]$ hadoop fs -cat /Hive_assignment/2019-Oct.csv | head
event_time,event_type,product_id,category_id,category_code,brand,price,user_id,user_session
2019-10-01 00:00:00 UTC,cart,5773203,1487580005134238553,,runail,2.62,463240011,26dd6e6e-4dac-4778-8d2c-92e149dab885
2019-10-01 00:00:03 UTC,cart,5773353,1487580005134238553,,runail,2.62,463240011,26dd6e6e-4dac-4778-8d2c-92e149dab885
2019-10-01 00:00:07 UTC,cart,5881589,2151191071051219817,,lovely,13.48,429681830,49e8d843-adf3-428b-a2c3-fe8bc6a307c9
2019-10-01 00:00:07 UTC,cart,5723490,1487580005134238553,,runail,2.62,463240011,26dd6e6e-4dac-4778-8d2c-92e149dab885
2019-10-01 00:00:15 UTC,cart,5881449,1487580013522845895,,lovely,0.56,429681830,49e8d843-adf3-428b-a2c3-fe8bc6a307c9
2019-10-01 00:00:16 UTC,cart,5857269,1487580005134238553,,runail,2.62,430174032,73dea1e7-664e-43f4-8b30-d32b9d5af04f
2019-10-01 00:00:19 UTC,cart,5739055,1487580008246412266,,kapous,4.75,377667011,81326ac6-daa4-4f0a-b488-fd0956a78733
2019-10-01 00:00:24 UTC,cart,5825598,1487580009445982239,,0.56,467916806,2f5b5546-b8cb-9ee7-7ecd-84276f8ef486
2019-10-01 00:00:25 UTC,cart,5698989,1487580006317032337,,1.27,385985999,d30965e8-1101-44ab-b45d-cc1bb9fae694
```

```
[hadoop@ip-172-31-91-114 ~]$ hadoop fs -cat /Hive_assignment/2019-Nov.csv | head
event_time,event_type,product_id,category_id,category_code,brand,price,user_id,user_session
2019-11-01 00:00:02 UTC,view,5802432,1487580009286598681,,0.32,562076640,09fafd6c-6c99-46b1-834f-33527f4de241
2019-11-01 00:00:09 UTC,cart,5844397,1487580006317032337,,2.38,553329724,2067216c-31b5-455d-alcc-af0575a34ffb
2019-11-01 00:00:10 UTC,view,5837166,1783999064103190764,,pnb,22.22,556138645,57ed222e-a54a-4907-9944-5a875c2d7f4f
2019-11-01 00:00:11 UTC,cart,5876812,1487580010100293687,,jessnail,3.16,564506666,186c1951-8052-4b37-adce-dd9644b1d5f7
2019-11-01 00:00:24 UTC,remove_from_cart,5826182,1487580007483048900,,3.33,553329724,2067216c-31b5-455d-alcc-af0575a34ffb
2019-11-01 00:00:24 UTC,remove_from_cart,5826182,1487580007483048900,,3.33,553329724,2067216c-31b5-455d-alcc-af0575a34ffb
2019-11-01 00:00:25 UTC,view,5856189,1487580009026551821,,runail,15.71,562076640,09fafd6c-6c99-46b1-834f-33527f4de241
2019-11-01 00:00:32 UTC,view,5837835,1933472286753424063,,3.49,514649199,432a4e95-375c-4b40-bd36-0fc039e77580
2019-11-01 00:00:34 UTC,remove_from_cart,5870838,1487580007675986893,,milv,0.79,429913900,2f0bfff3c-252f-4fe6-afcd-5d8a6a92839a
```

- After successfully adding data, it's time to Set the data in Hive
- Launch Hive

```
[hadoop@ip-172-31-91-114 ~]$ hive
Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j2.properties Async: false
hive> show databases;
OK
default
Time taken: 0.652 seconds, Fetched: 1 row(s)
hive>
```

- Creating database, creating tables

```
hive> create database if not exists Hive_casestudy;
OK
Time taken: 0.058 seconds
hive> describe database Hive_casestudy;
OK
hive_casestudy      hdfs://ip-172-31-91-114.ec2.internal:8020/user/hive/warehouse/hive_casestudy.db hadoop  USER
Time taken: 0.026 seconds, Fetched: 1 row(s)
hive> use Hive_casestudy;
OK
Time taken: 0.015 seconds
```

- Creating a table from the raw data by taking care of the data dictionary.

```
hive> CREATE EXTERNAL TABLE IF NOT EXISTS Ecom (
  > event_time timestamp,
  > event_type string,
  > product_id string,
  > category_id string,
  > category_code string,
  > brand string,
  > price float,
  > user_id bigint,
  > user_session string )
  > ROW FORMAT SERDE 'org.apache.hadoop.hive.serde2.OpenCSVSerde'
  > STORED AS TEXTFILE
  > LOCATION '/Hive_assignment'
  > TBLPROPERTIES ("skip.header.line.count"="1");
OK
Time taken: 0.058 seconds
```

```
hive> DESCRIBE Ecom;
OK
event_time          string          from deserializer
event_type          string          from deserializer
product_id          string          from deserializer
category_id         string          from deserializer
category_code       string          from deserializer
brand               string          from deserializer
price               string          from deserializer
user_id             string          from deserializer
user_session        string          from deserializer
Time taken: 0.078 seconds, Fetched: 9 row(s)
```

- Loading data from both files into this table.

```
hive> LOAD DATA INPATH '/Hive_assignment/2019-Oct.csv' into table Ecom;
Loading data to table hive_casestudy.ecom
OK
Time taken: 1.786 seconds
hive> LOAD DATA INPATH '/Hive_assignment/2019-Nov.csv' into table Ecom;
Loading data to table hive_casestudy.ecom
OK
Time taken: 0.566 seconds
```

- Checking the data after loading them into tables

```
hive> select * from Ecom limit 3;
OK
2019-11-01 00:00:02 UTC view      5802432 1487580009286598681      0.32    562076640      09fafd6c-6c99-46b1
-834f-33527f4de241
2019-11-01 00:00:09 UTC cart      5844397 1487580006317032337      2.38    553329724      2067216c-31b5-455d
-alcc-af0575a34ffb
2019-11-01 00:00:10 UTC view      5837166 1783999064103190764      pnb     22.22    556138645      57ed222e-a54a-4907
-9944-5a875c2d7f4f
Time taken: 1.887 seconds, Fetched: 3 row(s)
hive> select * from Ecom where month(cast(replace(event_time,'UTC','') as timestamp))=10 limit 3;
OK
2019-10-01 00:00:00 UTC cart      5773203 1487580005134238553      runail  2.62     463240011      26dd6e6e-4dac-4778
-8d2c-92e149dab885
2019-10-01 00:00:03 UTC cart      5773353 1487580005134238553      runail  2.62     463240011      26dd6e6e-4dac-4778
-8d2c-92e149dab885
2019-10-01 00:00:07 UTC cart      5881589 2151191071051219817      lovely  13.48    429681830      49e8d843-adf3-428b
-a2c3-fe8bc6a307c9
Time taken: 0.648 seconds, Fetched: 3 row(s)
```

- Creating table for data analysis with data in proper format.

```
hive> CREATE EXTERNAL TABLE IF NOT EXISTS Ecom_data (
> event_time timestamp,
> event_type string,
> product_id string,
> category_id string,
> category_code string,
> brand string,
> price float,
> user_id bigint,
> user_session string )
> ROW FORMAT DELIMITED FIELDS TERMINATED BY ','
> LINES TERMINATED BY '\n'
> STORED AS TEXTFILE;
OK
Time taken: 0.116 seconds
hive> DESCRIBE Ecom_data;
OK
event_time      timestamp
event_type      string
product_id      string
category_id     string
category_code   string
brand           string
price           float
user_id         bigint
user_session    string
Time taken: 0.046 seconds, Fetched: 9 row(s)
```

- Inserting data into this table.

```
hive> INSERT INTO Ecom_data
> select
> cast(replace(event_time,'UTC','') as timestamp),
> event_type,
> product_id,
> category_id,
> category_code,
> brand,
> cast(price as float),
> cast(user_id as bigint),
> user_session
> from Ecom;
Query ID = hadoop_20210529191313_5a9a8902-41bb-4b47-854f-ae95f1de7f2b
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1622306138328_0004)

-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    2         2         0         0         0         0
-----
VERTICES: 01/01 [=====>] 100% ELAPSED TIME: 117.35 s
-----
Loading data to table hive_casestudy.ecom_data
OK
Time taken: 127.611 seconds
```

Q: Find the total revenue generated due to purchases made in October. (using the Ecom_data table)

```
hive> select sum(price)
> from Ecom_data
> where month(event_time)=10 and event_type='purchase';
Query ID = hadoop_20210530131436_5cd8b346-a5eb-4d9f-8382-8d1b04ad2712
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1622369673465_0007)

Map 1: 0/7      Reducer 2: 0/1
Map 1: 0/7      Reducer 2: 0/1
Map 1: 0/7      Reducer 2: 0/1
Map 1: 0(+1)/7  Reducer 2: 0/1
Map 1: 0(+3)/7  Reducer 2: 0/1
Map 1: 0(+3)/7  Reducer 2: 0/1
Map 1: 0(+3)/7  Reducer 2: 0/1
Map 1: 0(+3)/7  Reducer 2: 0/1
Map 1: 0(+3)/7  Reducer 2: 0/1
Map 1: 1(+2)/7  Reducer 2: 0/1
Map 1: 1(+3)/7  Reducer 2: 0/1
Map 1: 2(+3)/7  Reducer 2: 0/1
Map 1: 3(+3)/7  Reducer 2: 0/1
Map 1: 3(+3)/7  Reducer 2: 0/1
Map 1: 5(+2)/7  Reducer 2: 0/1
Map 1: 5(+2)/7  Reducer 2: 0(+1)/1
Map 1: 6(+1)/7  Reducer 2: 0(+1)/1
Map 1: 7/7      Reducer 2: 0(+1)/1
Map 1: 7/7      Reducer 2: 1/1
OK
1211538.4295325726
Time taken: 35.613 seconds, Fetched: 1 row(s)
```

```
SELECT SUM(price)
```

```
FROM ecomm_data
```

```
WHERE MONTH(event_time)=10 AND event_type='purchase';
```

Time taken= 35 seconds.

Once the base table is created, we need to optimize the table for quick query result through partitioning and bucketing. Our optimized table name is Ecom_data_part.

Now we will Enabling Dynamic Partitioning and creating a partitioned table with buckets.

```
hive> set hive.exec.dynamic.partition=true;
hive> set hive.exec.dynamic.partition.mode=nonstrict;
hive> CREATE EXTERNAL TABLE IF NOT EXISTS Ecom_data_part (
> event_time timestamp,
> event_type string,
> product_id string,
> category_id string,
> category_code string,
> brand string,
> price float,
> user_id bigint,
> user_session string )
> PARTITIONED BY (year int, month int)
> CLUSTERED BY (category_id) into 4 buckets
> ROW FORMAT DELIMITED FIELDS TERMINATED BY ','
> LINES TERMINATED BY '\n'
> STORED AS TEXTFILE;
OK
Time taken: 0.115 seconds
hive> show tables;
OK
ecom
ecom_data
ecom_data_part
Time taken: 0.033 seconds, Fetched: 3 row(s)
```



```
hive> describe Ecom_data_part;
OK
event_time          timestamp
event_type          string
product_id          string
category_id         string
category_code       string
brand               string
price               float
user_id             bigint
user_session        string
year                int
month               int

# Partition Information
# col_name          data_type          comment
year                int
month               int
```

The new optimised table has been created. We will now insert the data into this table.

```
hive> INSERT INTO table Ecom_data_part PARTITION (year, month)
> select
> cast(replace(event_time,'UTC','') as timestamp),
> event_type,
> product_id,
> category_id,
> category_code,
> brand,
> cast(price as float),
> cast(user_id as bigint),
> user_session,
> year(cast(replace(event_time,'UTC','') as timestamp)),
> month(cast(replace(event_time,'UTC','') as timestamp))
> from Ecom;
Query ID = hadoop_20210530132016_22e509b9-26f6-4b97-8643-03374bfc66b
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1622369673465_0007)

-----
VERTICES      MODE           STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED   2       2          0        0        0       0
Reducer 2 ..... container  SUCCEEDED   5       5          0        0        0       0
-----
VERTICES: 02/02  [=====>>] 100%  ELAPSED TIME: 217.93 s
-----
Loading data to table default.ecom_data_part partition (year=null, month=null)

Loaded : 2/2 partitions.
Time taken to load dynamic partitions: 0.222 seconds
Time taken for adding to write entity : 0.0 seconds
OK
```

We will now test this table by running the same query in this optimised table and note the time taken.

```
hive> select sum(price)
> from Ecom data part
> where month(event_time)=10 and event_type='purchase';
Query ID = hadoop_20210530132728_7dc057d7-f76e-40cd-b949-b92e154e57ba
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1622369673465_0007)

-----
VERTICES      MODE           STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED   8       8          0        0        0       0
Reducer 2 ..... container  SUCCEEDED   1       1          0        0        0       0
-----
VERTICES: 02/02  [=====>>] 100%  ELAPSED TIME: 33.20 s
-----
OK
1211538.4295325726
Time taken: 33.966 seconds, Fetched: 1 row(s)
```

```
SELECT SUM(price) as total_revenue
FROM ecomm_data_part
```

```
WHERE MONTH(event_time)=10 AND event_type='purchase';
```

Time taken= 33.96 seconds.

Enabling Second Approach for Dynamic Partitioning and creating a partitioned table with buckets.

```
hive> set hive.exec.dynamic.partition=true;  
hive> set hive.exec.dynamic.partition.mode=nonstrict;
```

```
hive> CREATE EXTERNAL TABLE IF NOT EXISTS Ecom_data_part2 (  
  > event_time timestamp,  
  > product_id string,  
  > category_id string,  
  > category_code string,  
  > brand string,  
  > price float,  
  > user_id bigint,  
  > user_session string )  
  > PARTITIONED BY (event_type string)  
  > CLUSTERED BY (category_id) into 5 buckets  
  > ROW FORMAT SERDE 'org.apache.hadoop.hive.serde2.OpenCSVSerde'  
  > STORED as TEXTFILE;  
OK  
Time taken: 0.213 seconds
```

```
hive> describe Ecom_data_part2;  
OK  
event_time          string          from deserializer  
product_id          string          from deserializer  
category_id         string          from deserializer  
category_code       string          from deserializer  
brand               string          from deserializer  
price               string          from deserializer  
user_id             string          from deserializer  
user_session        string          from deserializer  
event_type          string  
  
# Partition Information  
# col_name           data_type       comment  
  
event_type           string  
Time taken: 0.1 seconds, Fetched: 14 row(s)
```

The new optimised table has been created. We will now insert the data into this table.

```
hive> INSERT INTO table Ecom_data_part2 PARTITION (event_type)  
  > select  
  > cast(replace(event_time,'UTC','') as timestamp),  
  > product_id,  
  > category_id,  
  > category_code,  
  > brand,  
  > cast(price as float),  
  > cast(user_id as bigint),  
  > user_session,  
  > event_type  
  > from Ecom;  
Query ID = hadoop_20210530152438_5b919786-9edc-420c-a8c5-a140974509df  
Total jobs = 1  
Launching Job 1 out of 1  
Tez session was closed. Reopening...  
Session re-established.  
Status: Running (Executing on YARN cluster with App id application_1622369673465_0015)  
  
-----  
VERTICES    MODE           STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED  
-----  
Map 1 ..... container    SUCCEEDED    2         2         0         0         0         0  
Reducer 2 ..... container    SUCCEEDED    5         5         0         0         0         0  
-----  
VERTICES: 02/02  [=====>>>] 100%  ELAPSED TIME: 213.76 s  
-----  
Loading data to table hive_casestudy.ecom_data_part2 partition (event_type=null)
```

```
-----  
VERTICES: 02/02 [=====>>] 100% ELAPSED TIME: 213.76 s  
-----  
Loading data to table hive_casestudy.ecom_data_part2 partition (event_type=null)  
  
Loaded : 4/4 partitions.  
Time taken to load dynamic partitions: 0.714 seconds  
Time taken for adding to write entity : 0.002 seconds  
OK  
Time taken: 223.6 seconds
```

We will now test this table by running the same query in this optimised table and note the time taken.

```
hive> select sum(price)  
> from Ecom_data_part2  
> where month(event_time)=10 and event_type='purchase';  
Query ID = hadoop_20210530153142_e05f26b5-dd66-4f47-94eb-1a96f6e39619  
Total jobs = 1  
Launching Job 1 out of 1  
Status: Running (Executing on YARN cluster with App id application_1622369673465_0015)  
  
-----  
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED  
-----  
Map 1 ..... container  SUCCEEDED    3         3         0         0         0         0  
Reducer 2 ..... container  SUCCEEDED    1         1         0         0         0         0  
-----  
VERTICES: 02/02 [=====>>] 100% ELAPSED TIME: 24.08 s  
-----  
OK  
1211538.4299998898  
Time taken: 26.316 seconds, Fetched: 1 row(s)
```

```
SELECT SUM(price) as total_revenue  
FROM ecomm_data_part2  
WHERE MONTH(event_time)=10 AND event_type='purchase';
```

Time taken= 26.316 seconds.

- So here we find that By **Partition by** over 'event_type' and **clustering by** 'category_id' we get the most optimized output of the query.

Q1. Find the total revenue generated due to purchases made in October.

```
hive> select sum(price) from ecom_data_part2 where month(event_time)=10 and event_type='purchase';  
Query ID = hadoop_20210530173820_5f776e91-c66b-4747-a414-6f967baef6e5  
Total jobs = 1  
Launching Job 1 out of 1  
Status: Running (Executing on YARN cluster with App id application_1622369673465_0016)  
  
-----  
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED  
-----  
Map 1 ..... container  SUCCEEDED    3         3         0         0         0         0  
Reducer 2 ..... container  SUCCEEDED    1         1         0         0         0         0  
-----  
VERTICES: 02/02 [=====>>] 100% ELAPSED TIME: 24.49 s  
-----  
OK  
1211538.4299998898  
Time taken: 29.151 seconds, Fetched: 1 row(s)
```

Query: select sum(price) from ecom_data_part where month(event_time)=10 and event_type='purchase';

Note :- The screenshot of the same query from both the base table and the bucketed table. When compared the bucketed table takes less time to query the result than the base table. This is the use of partitioning and bucketing the data

Q2. Write a query to yield the total sum of purchases per month in a single output.

```
hive> select month(event_time), sum(price) from Ecom_data_part2 where year(event_time)=2019 and event_type='purchase' group by month(event_time);
Query ID = hadoop_20210530174222_15e2c515-e813-440c-8da9-0286eld3a716
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1622369673465_0016)
```

VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	container	SUCCEEDED	3	3	0	0	0	0
Reducer 2	container	SUCCEEDED	1	1	0	0	0	0

```
VERTICES: 02/02 [=====>>] 100% ELAPSED TIME: 25.25 s
OK
10      1211538.4299998898
11      1531016.9
Time taken: 26.1 seconds, Fetched: 2 row(s)
```

Query: select month(event_time), sum(price) from Ecom_data_part2 where year(event_time)=2019 and event_type='purchase' group by month(event_time);

Q3. Write a query to find the change in revenue generated due to purchases from October to November.

```
hive> select sum (case when month(event_time)=10 then price else -1*price end) as change_in_revenue from ecom_data_part2 where month(event_time) in (10,11) and event_type='purchase';
Query ID = hadoop_20210530174823_652f7de3-ef3f-494f-ba14-36942e02bcc7
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1622369673465_0017)
```

VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	container	SUCCEEDED	3	3	0	0	0	0
Reducer 2	container	SUCCEEDED	1	1	0	0	0	0

```
VERTICES: 02/02 [=====>>] 100% ELAPSED TIME: 25.95 s
OK
-319478.47000012523
Time taken: 34.617 seconds, Fetched: 1 row(s)
```

Query: select sum (case when month(event_time)=10 then price else -1*price end) as change_in_revenue from ecom_data_part2 where month(event_time) in (10,11) and event_type='purchase';

Q4. Find distinct categories of products. Categories with null category code can be ignored.

Result: Total categories: 500 categories

```
hive> select distinct split(category_code,'\\\.')[0] as cat from Ecom_data_part2 where split(category_code,'\\\.')[0] <> '';
Query ID = hadoop_20210530180655_flaae807-f897-43f0-80d8-068c9ef675d3
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1622369673465_0019)

-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    6         6         0         0         0         0
Reducer 2 ..... container  SUCCEEDED    5         5         0         0         0         0
-----
VERTICES: 02/02  [=====>>>] 100%  ELAPSED TIME: 68.62 s
-----
OK
furniture
appliances
accessories
apparel
sport
stationery
Time taken: 76.878 seconds, Fetched: 6 row(s)
```

Query: select distinct split(category_code,'\\\.')[0] as cat from Ecom_data_part2 where split(category_code,'\\\.')[0] <> '';

Q5. Find the total number of products available under each category.

ans: 6 products available

```
hive> select split(category_code,'\\\.')[0] as cat, count(product_id) as no_of_products
> FROM Ecom_data_part2
> WHERE SPLIT(category_code,'\\\.')[0] <> ''
> GROUP BY SPLIT(category_code,'\\\.')[0]
> ORDER BY No_of_products DESC;
Query ID = hadoop_20210530181540_ea02db9c-8580-499e-9102-999c2bd0d932
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1622369673465_0019)

-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    6         6         0         0         0         0
Reducer 2 ..... container  SUCCEEDED    5         5         0         0         0         0
Reducer 3 ..... container  SUCCEEDED    1         1         0         0         0         0
-----
VERTICES: 03/03  [=====>>>] 100%  ELAPSED TIME: 67.61 s
-----
OK
appliances      61736
stationery      26722
furniture       23604
apparel 18232
accessories     12929
sport           2
Time taken: 68.327 seconds, Fetched: 6 row(s)
```

Query: SELECT SPLIT(category_code,'\\\.')[0] AS cat, COUNT(product_id) AS No_of_products FROM Ecom_data_part2 WHERE SPLIT(category_code,'\\\.')[0] <> '' GROUP BY SPLIT(category_code,'\\\.')[0] ORDER BY No_of_products DESC;

Q6. Which brand had the maximum sales in October and November combined?

```
hive> SELECT brand, SUM (price) AS sales FROM ecom_data_part2 WHERE BRAND <>' ' and event_type='purchase'
      GROUP BY brand ORDER BY sales DESC LIMIT 1
      > ;
Query ID = hadoop_20210530182225_d353f401-e44e-4127-afaa-a2cd4ee83fb9
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1622369673465_0019)

-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    3         3         0         0         0         0
Reducer 2 ..... container  SUCCEEDED    1         1         0         0         0         0
Reducer 3 ..... container  SUCCEEDED    1         1         0         0         0         0
-----
VERTICES: 03/03 [=====>>>] 100% ELAPSED TIME: 22.85 s
-----
OK
runail 148297.94000000049
Time taken: 23.46 seconds, Fetched: 1 row(s)
```

Query: SELECT brand, SUM (price) AS sales FROM ecom_data_part2 WHERE BRAND <>" and event_type='purchase' GROUP BY brand ORDER BY sales DESC LIMIT 1;

Q7. Which brands increased their sales from October to November?

```
hive> WITH Monthly_rev AS (
  > SELECT brand,
  > SUM(CASE WHEN date_format(event_time, 'MM')=10 THEN price ELSE 0 END) AS Oct_rev,
  > SUM(CASE WHEN date_format(event_time, 'MM')=11 THEN price ELSE 0 END) AS Nov_rev
  > FROM Ecom_data_part2
  > WHERE event_type='purchase' AND date_format(event_time, 'MM') IN ('10', '11')
  > GROUP BY brand )
  > SELECT brand, Oct_rev, Nov_rev, Nov_rev-Oct_rev AS Sales_diff
  > FROM Monthly_rev
  > WHERE (Nov_rev - Oct_rev)>0
  > ORDER BY Sales_diff;
Query ID = hadoop_20210530182619_e806a32f-33c7-4200-95c5-652bddd1cf17
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1622369673465_0019)

-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    3         3         0         0         0         0
Reducer 2 ..... container  SUCCEEDED    1         1         0         0         0         0
Reducer 3 ..... container  SUCCEEDED    1         1         0         0         0         0
-----
VERTICES: 03/03 [=====>>>] 100% ELAPSED TIME: 27.23 s
-----
OK
ovale 2.54 3.1 0.56
cosima 20.23 20.930000000000003 0.7000000000000028
grace 100.920000000000002 102.61000000000004 1.69000000000000261
helloganic 0.0 3.1 3.1
```

```
artex 2730.6399999999994 4327.249999999997 1596.6099999999979
beautix 10493.949999999998 12222.949999999997 1729.0000000000011
milv 3904.9399999999983 5642.0099999999976 1737.0699999999993
masura 31266.079999999823 33058.4699999998706 1792.3900000000476
f.o.x 6624.23 8577.280000000001 1953.05000000000102
kapous 11927.1600000000113 14093.0800000000078 2165.9199999999655
concept 11032.139999999974 13380.399999999994 2348.2599999999657
estel 21756.7500000000084 24142.670000000007 2385.9199999999873
kaypro 881.34 3268.7000000000003 2387.36
benovy 409.6199999999999 3259.969999999992 2850.349999999992
italwax 21940.23999999973 24799.369999999766 2859.1300000000374
yoko 8756.909999999994 11707.879999999965 2950.9699999999702
haruyama 9390.690000000014 12352.90999999999 2962.219999999985
marathon 7280.749999999997 10273.099999999986 2992.3499999999885
lovely 8704.37999999999 11939.0600000000029 3234.6800000000385
bpw.style 11572.1500000001808 14837.4400000002425 3265.29000000006175
staleks 8519.7300000000023 11875.609999999999 3355.8799999999756
freedecor 3421.7799999999706 7671.8000000000216 4250.0200000000245
runail 71539.279999999619 76758.659999999736 5219.3800000001169
polarus 6013.72000000000075 11371.9300000000013 5358.21000000000055
cosmoprofi 8322.809999999996 14536.989999999958 6214.179999999962
jessnail 26287.840000000013 33345.230000000008 7057.389999999952
strong 29196.629999999994 38671.269999999994 9474.64
ingarden 23161.390000000044 33566.20999999995 10404.8199999999057
lianail 5892.839999999998 16394.2400000000194 10501.4000000000214
uno 35302.030000000014 51039.7499999998894 15737.7199999998757
grattol 35445.540000000078 71472.70999999995 36027.169999999872
474679.059999999656 619509.23999999899 144830.17999999933
Time taken: 27.859 seconds, Fetched: 161 row(s)
```

```
Query: WITH Monthly_rev AS (
SELECT brand,
SUM(CASE WHEN date_format(event_time, 'MM')=10 THEN price ELSE 0 END) AS Oct_rev,
SUM(CASE WHEN date_format(event_time, 'MM')=11 THEN price ELSE 0 END) AS Nov_rev
FROM Ecom_data_part2
WHERE event_type='purchase' AND date_format(event_time, 'MM') IN ('10', '11')
GROUP BY brand )
SELECT brand, Oct_rev, Nov_rev, Nov_rev-Oct_rev AS Sales_diff
FROM Monthly_rev
WHERE (Nov_rev - Oct_rev)>0
ORDER BY Sales_diff;
```

So from the above data we can see that there are total of 161 brands increasing from October to November.

Q8. Your company wants to reward the top 10 users of its website with a Golden Customer plan. Write a query to generate a list of top 10 users who spend the most.

```
hive> SELECT user_id, SUM(price) as Total_Expense
> FROM Ecom_data_part2
> WHERE event_type='purchase'
> GROUP BY user_id
> ORDER BY Total_Expense DESC
> LIMIT 10;
Query ID = hadoop_20210530183207_02c14d10-5fad-4f6f-8e53-28303a610f2a
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1622369673465_0020)

-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED   3         3         0         0         0         0
Reducer 2 ..... container  SUCCEEDED   1         1         0         0         0         0
Reducer 3 ..... container  SUCCEEDED   1         1         0         0         0         0
-----
VERTICES: 03/03  [=====>>] 100%  ELAPSED TIME: 25.11 s
-----
OK
557790271      2715.8699999999995
150318419      1645.97
562167663      1352.85
531900924      1329.4499999999998
557850743      1295.48
522130011      1185.3899999999999
561592095      1109.7000000000003
-----

OK
557790271      2715.8699999999995
150318419      1645.97
562167663      1352.85
531900924      1329.4499999999998
557850743      1295.48
522130011      1185.3899999999999
561592095      1109.7000000000003
431950134      1097.5899999999997
566576008      1056.3600000000006
521347209      1040.91
Time taken: 33.116 seconds, Fetched: 10 row(s)
```

Query: SELECT user_id, SUM(price) as Total_Expense

FROM Ecom_data_part2

WHERE event_type='purchase'

GROUP BY user_id

ORDER BY Total_Expense DESC

LIMIT 10;